

NIHAR DUGADE

New York, NY | Open to relocation | +15513594372 | dugadenihar80@gmail.com | [LinkedIn](#) | [Website](#) | [Github](#)

PROFESSIONAL SUMMARY

Analytics Engineer with experience building scalable ETL/ELT pipelines, real-time streaming architecture, and data quality using AWS, Snowflake, Spark, and Kafka. Enhanced reporting speed and accuracy, and delivered reliable Kafka pipelines for marketplace data, ready to drive high-performance analytics for data-focused teams.

WORK EXPERIENCE

Analytics Engineer **Oct 2025 - Present**
Just Innovation Fund LLC *Remote, NY*

- Engineered PySpark ETL/ELT pipelines to unify product, finance, and usage data across JustLabs' voucher platform into a single reporting layer cutting reporting latency by 35% and enabling near real-time insights for faster pricing decisions
- Integrated automated PySpark and SQL data quality checks at the pipeline ingestion layer to catch bad voucher and transaction data before it reached executive dashboards improving reporting accuracy by 22% and cutting manual validation effort by 40%
- Architected dimensional models (fact and dimension tables) in **Snowflake** to support **cohort analytics** and buyer segment reporting across the voucher platform improving query performance by 30% and cutting analyst query time by 25%
- Standardized metric definitions for revenue, usage, and engagement across product, finance, and analytics teams by building shared dbt data models achieving 99%+ data availability and reducing reporting turnaround by 30%

Data Engineering Intern **Jan 2025 - May 2025**
Cantonica *Remote, NY*

- Built real-time **Kafka** event ingestion pipelines backed by MongoDB to reliably capture app interactions and user discovery events across Cantonica's PWA marketplace achieving 100% message delivery and eliminating production data loss
- Implemented **Spark** Streaming pipelines with checkpointing to keep real-time app listing and engagement analytics current even during pipeline failures maintaining 99.9% uptime and cutting recovery time by 40%
- Accelerated same-day app indexing for newly submitted PWAs by designing distributed PySpark ETL pipelines on Spark clusters reducing processing latency by ~70% and ensuring developers saw their apps live within hours of submission
- Streamlined pipeline deployments across dev, staging, and production by containerizing services with Docker and setting up CI/CD automation cutting release cycles by 30% and eliminating environment mismatch failures

Data Analyst **Jul 2021 - Sept 2022**
LRA Packaging *Faridabad, India*

- Modernized manual end-of-shift data exports with continuous AWS Glue, EMR, and S3 ETL pipelines that ingested packaging line data throughout the day reducing reporting latency by 40% and giving plant managers real-time production visibility
- Crafted a unified Snowflake star-schema data warehouse for packaging operations so every department accessed consistent plant performance metrics without building one-off queries cutting ad-hoc requests by 35% and improving query speed by 30%
- Developed real-time IoT sensor pipelines using **Kafka** on **AWS MSK** to detect early warning signs of machine failures on chip packaging lines reducing unplanned downtime by 20% and enabling proactive maintenance scheduling
- Enforced data quality rules at the ingestion layer using **AWS Glue Catalog** and Lake Formation to prevent bad sensor readings from flowing into production reports increasing data consistency by 30%

PROJECTS

NYC Taxi Medallion Lakehouse Pipeline

- Established a production-grade medallion pipeline on 67M NYC Taxi records using Delta Lake, DuckDB, and dbt, enabling scalable bronze-silver-gold processing with ACID transactions and incremental CDC handling.
- Automated data quality and observability across all pipeline layers using Great Expectations, Prefect, OpenLineage, and GitHub Actions CI/CD, catching schema violations and quality issues before downstream reporting.

Sentinel MCP Data Governance Agent

- Deployed a Python data governance agent using FastMCP, OpenAI, SQLAlchemy, and PostgreSQL to catch schema violations, null anomalies, and duplicate SKUs across dirty inventory tables via a controlled MCP bridge reducing manual audit time by ~60%
- Designed a ReAct agent workflow with SQLAlchemy audit tools to flag anomalies, profile nulls, and auto-generate remediation SQL into a structured Markdown report on every run

SKILLS

Languages: Python, SQL (Advanced), R, Scala, DAX

Data Engineering: ETL/ELT Pipelines, Data Modeling (Star Schema, Snowflake Schema), Data Warehousing, Data Quality, Data Governance, Pipeline Optimization, Parquet, Delta Lake, Apache Iceberg

Big Data & Streaming: Apache Spark (PySpark, Spark Streaming), Kafka, Apache Flink, Databricks, Real-Time & Batch Pipelines

Cloud Platforms & BI: AWS (S3, Glue, Lambda, EMR, MSK, Redshift, Athena, Lake Formation), Snowflake, BigQuery, Tableau, Power BI

Databases: PostgreSQL, MySQL, DynamoDB, MongoDB (NoSQL)

Orchestration & Transformation: Airflow, dbt, dbt Cloud, Prefect, MLflow, Pipeline Automation

AI & GenAI: LLM Integration, Prompt Engineering, MCP, Codex, LangChain, Claude Code, GitHub Copilot

DevOps: Docker, Git, CI/CD, CloudWatch, Linux

EDUCATION

Master of Science, Data Science | *Stevens Institute of Technology, Hoboken*

May 2025

Bachelor of Science, Physics | *DG Ruparel College, Mumbai, India*

May 2022