

NIHAR DUGADE

New York, NY | Open to relocation | +15513594372 | [Email](#) | [LinkedIn](#) | [Github](#)

WORK EXPERIENCE

Just Innovation Fund LLC

Remote, NY

Data Analyst

Oct 2025 – Present

- Engineered PySpark ETL/ELT pipelines to unify product, finance, and usage data across JustLabs' voucher platform, cutting reporting latency by 35% and enabling near real-time insights for faster pricing decisions
- Resolved **KPI inconsistencies** eroding executive confidence by building **PySpark** and **SQL** anomaly detection frameworks that caught bad voucher and transaction data at ingestion, improving dashboard accuracy by 22% and cutting manual checks by 40%
- Architected **Snowflake** and **Redshift** dimensional models (fact and dimension tables) to support cohort analytics across JustLabs' buyer segments and group purchasing behavior, improving query performance by 30% and cutting analyst query time by 25%
- Eliminated reporting bottlenecks slowing visibility into bulk order and redemption pipeline performance by refactoring large-scale **Snowflake** and **Redshift SQL** transformations through partition pruning and join optimization, cutting execution time by 45%
- Standardized conflicting metric definitions across product, finance, and analytics by translating ambiguous stakeholder requirements into shared production-grade SQL models, achieving 99%+ data availability and reducing delivery turnaround by 30%

Cantonica

Remote, NY

Data Engineer Intern

Jan 2025 – May 2025

- Designed fault-tolerant **Kafka** and **MongoDB** event pipelines with schema evolution to capture every user interaction on **Cantonica's PWA** marketplace without loss, achieving 100% message delivery and reducing data loss incidents to near zero
- Implemented **Spark Streaming** pipelines with checkpointing and failure recovery so Cantonica's real-time app listing and engagement analytics never went stale during incidents, maintaining **99.9% uptime** and cutting pipeline recovery time by **40%**
- Accelerated same-day app indexing for newly submitted PWAs by designing distributed **PySpark ETL** pipelines on **Spark clusters**, reducing processing latency by **~70%** and ensuring developers saw their apps live on the marketplace within hours of submission
- Streamlined **deployment** of data pipelines across Cantonica's marketplace by containerizing services with **Docker** and establishing **CI/CD** workflows, cutting release cycles by 30% and eliminating environment mismatch failures between dev and production
- Collaborated with engineering and data science teams to align pipeline architecture around shared app usage and engagement data standards, improving data accessibility and reducing ad-hoc requests by 25% across both teams

LRA Packaging

Remote, India

Data Analyst

Jul 2021 – Sep 2022

- Modernized manual shift-end reporting by implementing cloud-native **AWS Glue**, **EMR**, **Lambda**, and **S3 ETL** pipelines to continuously ingest LRA's packaging line data, reducing latency by 40% and giving plant managers live production visibility
- Eliminated conflicting department reports by designing **Snowflake** and **Redshift** star-schema models that gave every team a single source of truth for plant output metrics, cutting ad-hoc requests by 35% and improving query speed by 30%
- Enabled proactive maintenance on chip packaging lines by building real-time **Kafka** on **AWS MSK IoT** pipelines to surface machine failure patterns before stoppages occurred, reducing operational downtime by 20%
- Prevented bad sensor readings and operational data from corrupting production reports by building **AWS Glue Catalog** and **Lake Formation** governance frameworks with validation rules enforced at ingestion, increasing data consistency by **30%**
- Ensured shift-end reports reached plant managers before handovers by tuning complex **multi-table SQL** queries across production, machine, and logistics data through targeted indexing, partitioning, and join optimization, improving query response time by **25%**

SKILLS

Languages: Python, SQL (Advanced), R, Scala, C++, Go

Data Engineering: ETL/ELT Pipelines, Data Modeling (Star Schema, Snowflake), Data Warehousing, Data Integration, Data Quality, Data Governance, Pipeline Optimization

Big Data & Streaming: Apache Spark (PySpark, Spark Streaming), Kafka, Distributed Processing, Real-Time & Batch Pipelines

Cloud Platforms & Reporting: AWS (S3, Glue, Lambda, EMR, MSK, Redshift, Athena, Lake Formation), Data Lakes, Tableau

Databases: PostgreSQL, MySQL, DynamoDB, MongoDB (NoSQL), Data Warehouse Systems

Orchestration & Transformation: Airflow, dbt, Workflow Scheduling, Pipeline Automation

AI Tools & DevOps: Claude code, Gemini enterprise, Github copilot, dbt Cloud, Docker, Git, CI/CD, CloudWatch, Linux

PROJECTS

Retail Sales & AAPL Forecasting Pipeline

- Designed an end-to-end **time-series forecasting** system using **Python (ARIMA/SARIMA)**, automating data ingestion, transformation, and validation to reduce preprocessing time by 40% and accelerate forecasting cycles
- Standardized reusable workflows with ADF and Ljung-Box validation checks to ensure forecast-ready data quality, improving prediction reliability across scalable time-series analytics

HR Attrition Prediction System

- Built a Python ML pipeline for feature engineering and model training to predict employee churn, achieving ROC-AUC of 0.83
- Detected high-risk employee segments through analysis and model tuning, driving insights that reduced attrition by **~5%**

EDUCATION

Master of Science, Data Science | Stevens Institute of Technology, Hoboken

May 2025

Bachelor of Science, Physics | DG Ruparel College, Mumbai, India

June 2022